

1. Importing all the necessary dependencies

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
import plotly.io as plio
plio.templates
import plotly.express as px
import plotly.graph_objects as go

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

from sklearn.linear_model import LinearRegression, Ridge, Lasso, ElasticNet
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor

from xgboost import XGBRegressor

import joblib

from warnings import filterwarnings
filterwarnings(action='ignore')
```

2. Exploring the dataset

In [2]:

```
data = pd.read_csv('car data.csv')
data
```

Out[2]:

	Car_Name	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	Transmission
0	ritz	2014	3.35	5.59	27000	Petrol	Dealer	M
1	sx4	2013	4.75	9.54	43000	Diesel	Dealer	M
2	ciaz	2017	7.25	9.85	6900	Petrol	Dealer	M
3	wagon r	2011	2.85	4.15	5200	Petrol	Dealer	M
4	swift	2014	4.60	6.87	42450	Diesel	Dealer	M
...
296	city	2016	9.50	11.60	33988	Diesel	Dealer	M
297	brio	2015	4.00	5.90	60000	Petrol	Dealer	M
298	city	2009	3.35	11.00	87934	Petrol	Dealer	M
299	city	2017	11.50	12.50	9000	Diesel	Dealer	M
300	brio	2016	5.30	5.90	5464	Petrol	Dealer	M

301 rows × 9 columns

In [3]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Car_Name         301 non-null    object  
 1   Year             301 non-null    int64  
 2   Selling_Price    301 non-null    float64 
 3   Present_Price   301 non-null    float64 
 4   Driven_kms      301 non-null    int64  
 5   Fuel_Type        301 non-null    object  
 6   Selling_type     301 non-null    object  
 7   Transmission     301 non-null    object  
 8   Owner            301 non-null    int64  
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

Finding any duplicates if present

```
In [4]: duplicate_rows = data[data.duplicated(keep=False)]
duplicate_rows
```

Out[4]:

	Car_Name	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	Transmis
15	ertiga	2016	7.75	10.79	43000	Diesel	Dealer	Ma
17	ertiga	2016	7.75	10.79	43000	Diesel	Dealer	Ma
51	fortuner	2015	23.00	30.61	40000	Diesel	Dealer	Auto
93	fortuner	2015	23.00	30.61	40000	Diesel	Dealer	Auto

```
In [5]: duplicated_rows = data[data.duplicated(keep='first')]
duplicated_rows
```

Out[5]:

	Car_Name	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type	Selling_type	Transmis
17	ertiga	2016	7.75	10.79	43000	Diesel	Dealer	Ma
93	fortuner	2015	23.00	30.61	40000	Diesel	Dealer	Auto

Removing the duplicate rows from the dataset

```
In [6]: data_cleaned = data.drop(duplicated_rows.index)
#data_cleaned.to_csv('car_data_cleaned.csv', index=False)
```

```
In [7]: data = pd.read_csv('car_data_cleaned.csv')
```

```
In [8]: data.Year.value_counts()
```

Out[8]:

2015	60
2016	49
2014	38
2017	35
2013	33
2012	23
2011	19
2010	15
2008	7
2009	6
2006	4
2005	4
2003	2
2007	2
2018	1
2004	1

Name: Year, dtype: int64

```
In [9]: year = data.Year.value_counts().keys()
no_cars_sold = data.Year.value_counts().values
labels = [str(val) for val in no_cars_sold]

fig = px.bar(x=year, y=no_cars_sold, text=labels, title='Year vs No. of Cars Sold')

# Set the Labels for the y-axis and x-axis
fig.update_yaxes(title_text='no_cars_sold')
fig.update_xaxes(title_text='Year')

# Show the plot
fig.show()
```



```
In [10]: car_name = data.Car_Name.value_counts().head(25).keys()
cars_sold = data.Car_Name.value_counts().head(25).values
labels = [str(val) for val in cars_sold]

fig = px.bar(x=car_name, y=cars_sold, text=labels, title='Top 25 Car Name vs No. of Cars Sold')

# Set the Labels for the y-axis and x-axis
fig.update_yaxes(title_text='No. of Cars Sold')
fig.update_xaxes(title_text='Car Name')

# Show the plot
fig.show()
```

```
In [11]: data.Fuel_Type.value_counts().head(25)
```

```
Out[11]: Petrol    239
Diesel     58
CNG        2
Name: Fuel_Type, dtype: int64
```

```
In [12]: fuel_type = data.Fuel_Type.value_counts().keys()
cars_count = data.Fuel_Type.value_counts().values
labels = [str(val) for val in cars_count]

fig = px.bar(x=fuel_type, y=cars_count, text=labels, title='Fuel Type vs Cars')

# Set the Labels for the y-axis and x-axis
fig.update_yaxes(title_text='No. of Cars')
fig.update_xaxes(title_text='Fuel Type')

# Show the plot
fig.show()
```

```
In [13]: selling_type = data.Selling_type.value_counts().keys()
cars_count = data.Selling_type.value_counts().values
labels = [str(val) for val in cars_count]

fig = px.bar(x=selling_type, y=cars_count, text=labels, title='Selling Type vs

# Set the Labels for the y-axis and x-axis
fig.update_yaxes(title_text='No. of Cars')
fig.update_xaxes(title_text='Selling Type')

# Show the plot
fig.show()
```

```
In [14]: transmission_type = data.Transmission.value_counts().keys()
cars_count = data.Transmission.value_counts().values
labels = [str(val) for val in cars_count]

fig = px.bar(x=transmission_type, y=cars_count, text=labels, title='Transmission Type')

# Set the Labels for the y-axis and x-axis
fig.update_yaxes(title_text='No. of Cars')
fig.update_xaxes(title_text='Transmission Type')

# Show the plot
fig.show()
```

```
In [15]: ownership_type = data.Owner.value_counts().keys()
cars_count = data.Owner.value_counts().values
labels = [str(val) for val in cars_count]

fig = px.bar(x=ownership_type, y=cars_count, text=labels, title='Ownership Type')

# Set the Labels for the y-axis and x-axis
fig.update_yaxes(title_text='No. of Cars')
fig.update_xaxes(title_text='Ownership Type')

# Show the plot
fig.show()
```

Changing the categorical columns Fuel_Type, Selling_Type, Transmission and Owner Type to Numerical columns

```
In [16]: # Changing Fuel Type column : Petrol -> 0, Diesel -> 1 and CNG -> 2
category_mapping = {'Petrol': 0, 'Diesel': 1, 'CNG': 2}
data['Fuel_Type_encoded'] = data['Fuel_Type'].map(category_mapping)
```

```
In [17]: data.Fuel_Type_encoded.value_counts()
```

```
Out[17]: 0    239  
1     58  
2      2  
Name: Fuel_Type_encoded, dtype: int64
```

```
In [18]: fuel_type = data.Fuel_Type_encoded.value_counts().keys()  
cars_count = data.Fuel_Type_encoded.value_counts().values  
labels = [str(val) for val in cars_count]  
  
fig = px.bar(x=fuel_type, y=cars_count, text=labels, title='Fuel Type(encoded)  
  
# Set the labels for the y-axis and x-axis  
fig.update_yaxes(title_text='No. of Cars')  
fig.update_xaxes(title_text='Fuel Type(encoded)')  
  
# Show the plot  
fig.show()
```

```
In [19]: # Changing Selling_Type column : Dealer -> 0, Individual -> 1
category_mapping = {'Dealer': 0, 'Individual': 1}
data['Selling_type_encoded'] = data['Selling_type'].map(category_mapping)
```

```
In [20]: data.Selling_type_encoded.value_counts()
```

```
Out[20]: 0    193
1    106
Name: Selling_type_encoded, dtype: int64
```

```
In [21]: selling_type = data.Selling_type_encoded.value_counts().keys()
cars_count = data.Selling_type_encoded.value_counts().values
labels = [str(val) for val in cars_count]

fig = px.bar(x=selling_type, y=cars_count, text=labels, title='Selling Type(encoded)')

# Set the labels for the y-axis and x-axis
fig.update_yaxes(title_text='No. of Cars(encoded)')
fig.update_xaxes(title_text='Selling Type')

# Show the plot
fig.show()
```

```
In [22]: # Changing Transmission column : Manual -> 0, Automatic -> 1
category_mapping = {'Manual': 0, 'Automatic': 1}
data['Transmission_encoded'] = data['Transmission'].map(category_mapping)
```

```
In [23]: data.Transmission_encoded.value_counts()
```

```
Out[23]: 0    260
1     39
Name: Transmission_encoded, dtype: int64
```

```
In [24]: transmission_type = data.Transmission_encoded.value_counts().keys()
cars_count = data.Transmission_encoded.value_counts().values
labels = [str(val) for val in cars_count]

fig = px.bar(x=transmission_type, y=cars_count, text=labels, title='Transmission Type')

# Set the labels for the y-axis and x-axis
fig.update_yaxes(title_text='No. of Cars')
fig.update_xaxes(title_text='Transmission Type(encoded)')

# Show the plot
fig.show()
```

```
In [25]: # Changing Owner Type column : 0 -> 0, 1 -> 1 and 3 -> 2
category_mapping = {0: 0, 1: 1, 3: 2}
data['Owner_encoded'] = data['Owner'].map(category_mapping)
```

```
In [26]: data.Owner_encoded.value_counts()
```

```
Out[26]: 0    288
1     10
2      1
Name: Owner_encoded, dtype: int64
```

```
In [27]: ownership_type = data.Owner_encoded.value_counts().keys()
cars_count = data.Owner_encoded.value_counts().values
labels = [str(val) for val in cars_count]

fig = px.bar(x=ownership_type, y=cars_count, text=labels, title='Ownership Type')

# Set the labels for the y-axis and x-axis
fig.update_yaxes(title_text='No. of Cars')
fig.update_xaxes(title_text='Ownership Type(encoded)')

# Show the plot
fig.show()
```

```
In [28]: selected_columns = ['Car_Name', 'Year', 'Selling_Price', 'Present_Price', 'Driven_kms', 'Fuel_Type_encoded', 'Selling_type']
```

```
In [29]: new_data = data[selected_columns]
```

```
In [30]: new_data
#new_data.to_csv('cars price encoded.csv', index=False)
```

```
Out[30]:
```

	Car_Name	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type_encoded	Selling_type
0	ritz	2014	3.35	5.59	27000		0
1	sx4	2013	4.75	9.54	43000		1
2	ciaz	2017	7.25	9.85	6900		0
3	wagon r	2011	2.85	4.15	5200		0
4	swift	2014	4.60	6.87	42450		1
...
294	city	2016	9.50	11.60	33988		1
295	brio	2015	4.00	5.90	60000		0
296	city	2009	3.35	11.00	87934		0
297	city	2017	11.50	12.50	9000		1
298	brio	2016	5.30	5.90	5464		0

299 rows × 9 columns

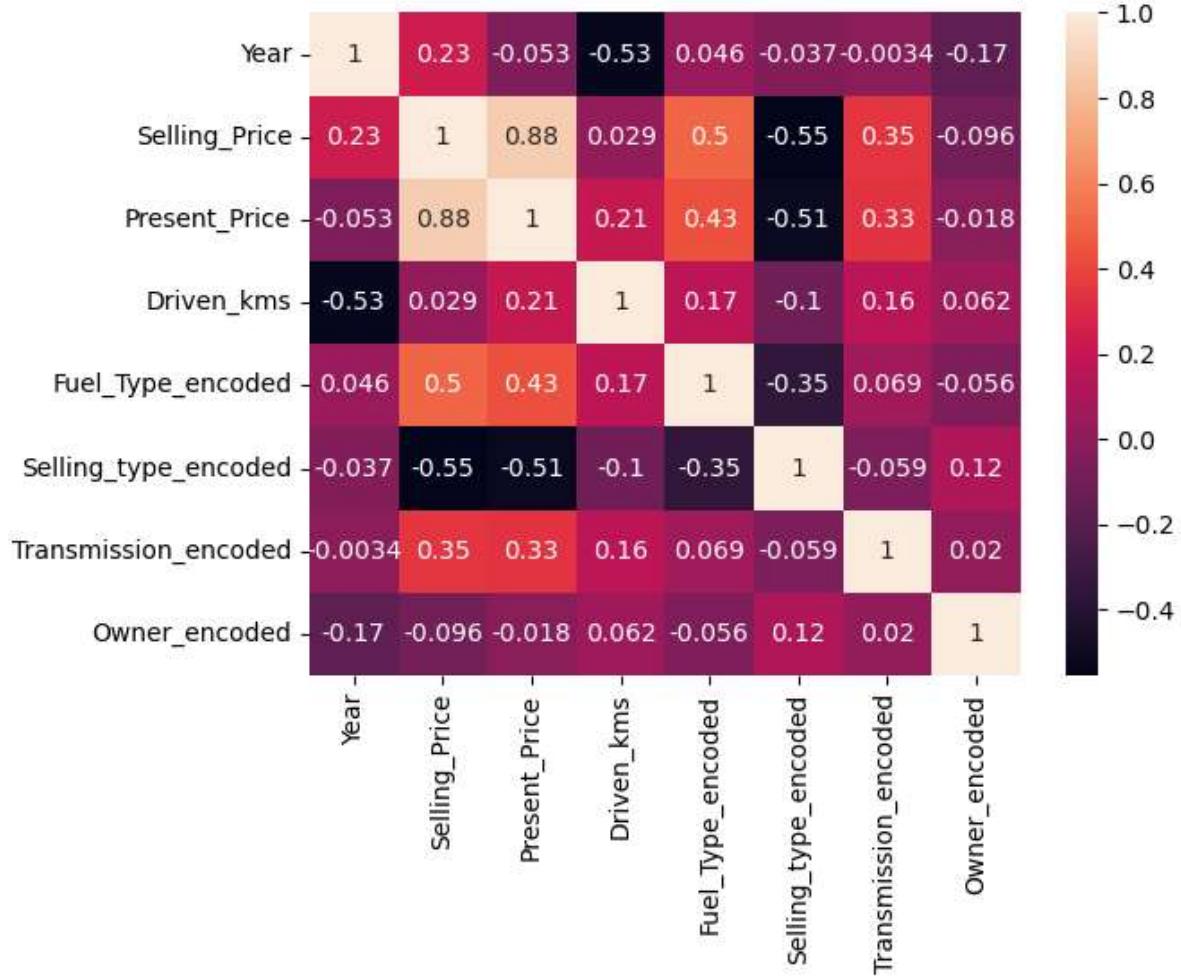
```
In [31]: new_data.iloc[:,1:].corr()
```

```
Out[31]:
```

	Year	Selling_Price	Present_Price	Driven_kms	Fuel_Type_encoded	Selling_type
Year	1.000000	0.234369	-0.053167	-0.525714	0.046210	
Selling_Price	0.234369	1.000000	0.876305	0.028566	0.500292	
Present_Price	-0.053167	0.876305	1.000000	0.205224	0.431887	
Driven_kms	-0.525714	0.028566	0.205224	1.000000	0.167287	
Fuel_Type_encoded	0.046210	0.500292	0.431887	0.167287	1.000000	
Selling_type_encoded	-0.036820	-0.553851	-0.511779	-0.101030	-0.347922	
Transmission_encoded	-0.003434	0.348869	0.334326	0.163881	0.068618	
Owner_encoded	-0.170694	-0.096019	-0.018158	0.061924	-0.055526	

In [32]: `sns.heatmap(new_data.iloc[:,1:].corr(), annot=True)`

Out[32]: <AxesSubplot:>



From the above correlation plot we can infer the following :-

A -> Selling_Price is influenced by :

(i) Present_price by a score of 0.88

(ii) Fuel_Type_encoded by a score of 0.5

(iii) Transmission_type by a score of 0.35

B -> Present_Price is influenced by :

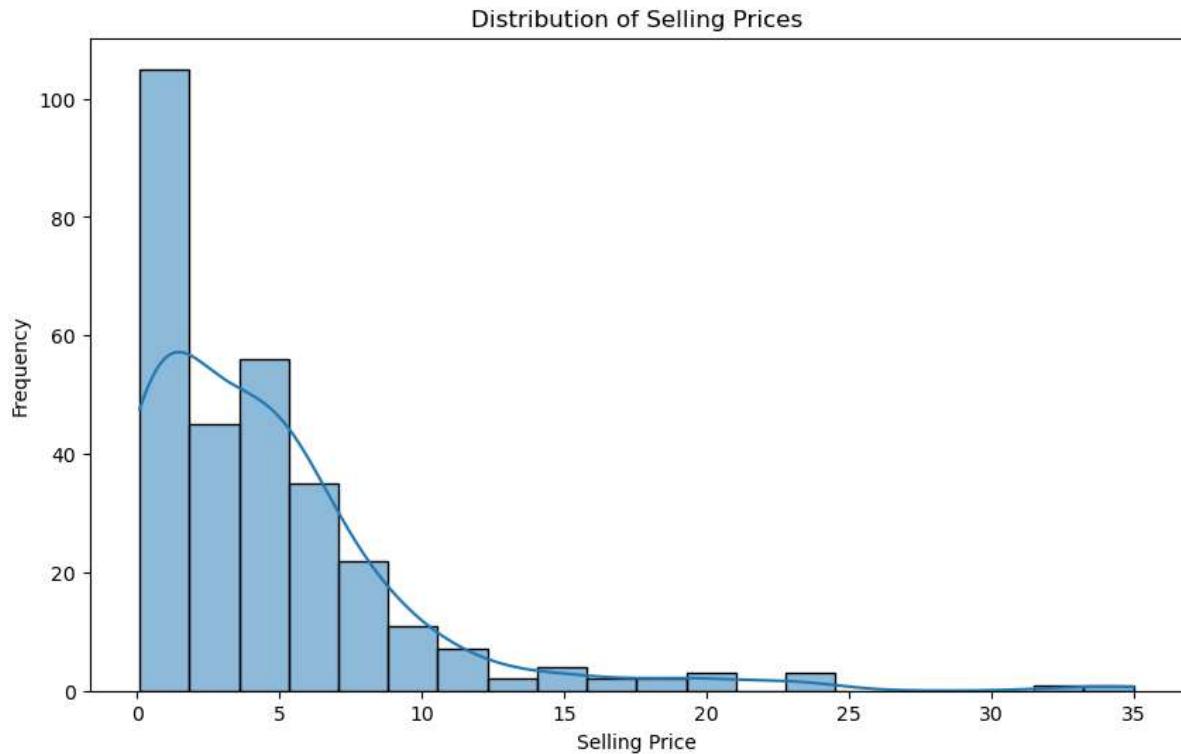
(i) Selling_price by a score of 0.88

(ii) Fuel_Type_encoded by a score of 0.43

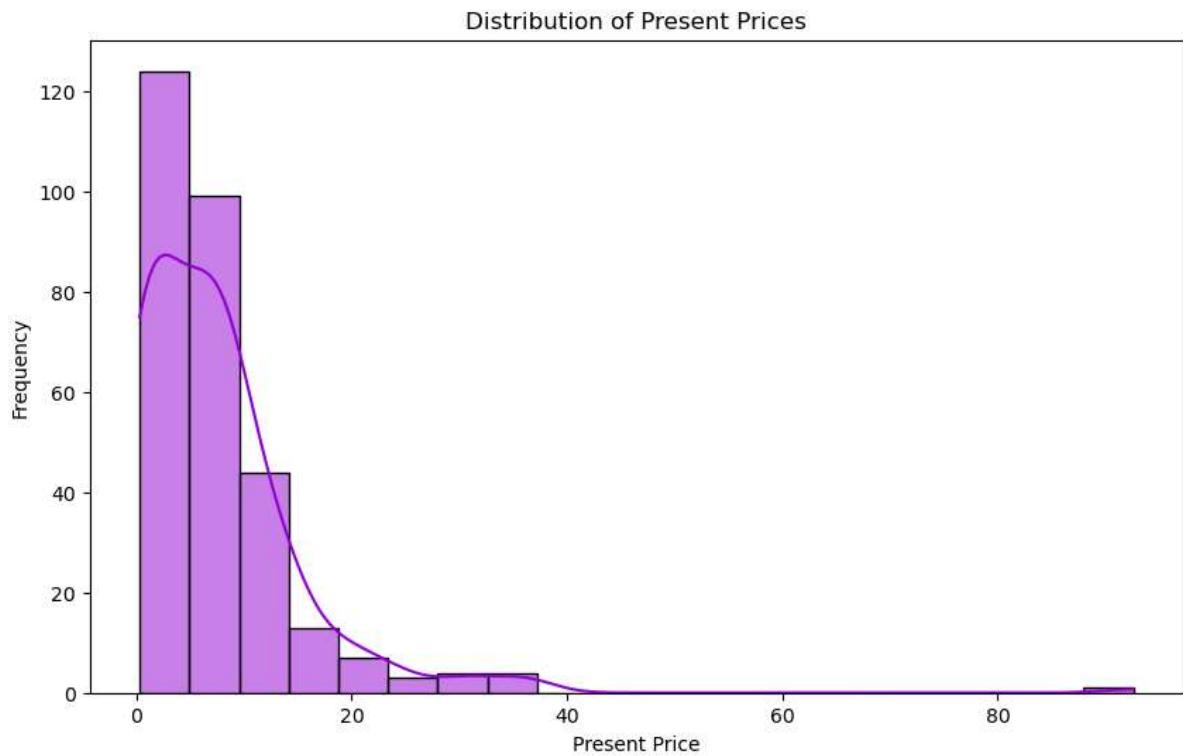
(iii) Transmission_type by a score of 0.33

3. Analyzing the distribution using histogram plots

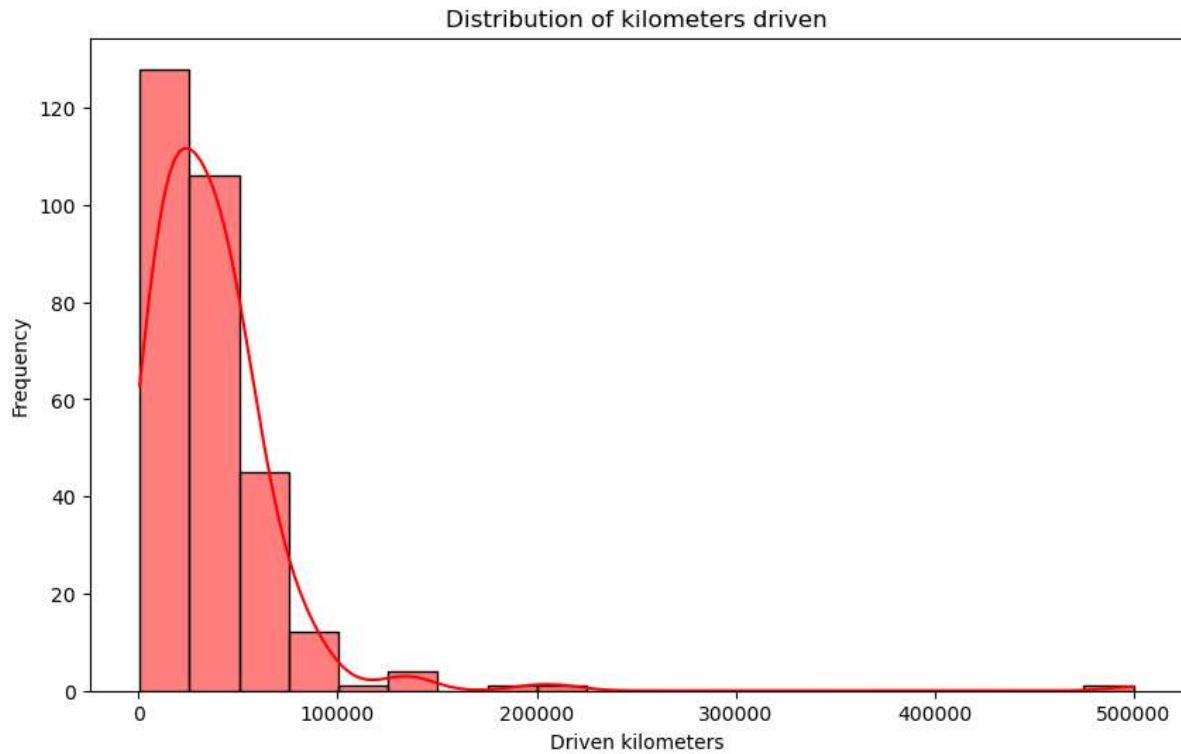
```
In [33]: plt.figure(figsize=(10, 6))
sns.histplot(new_data['Selling_Price'], bins=20, kde=True)
plt.xlabel('Selling Price')
plt.ylabel('Frequency')
plt.title('Distribution of Selling Prices')
plt.show()
```



```
In [34]: plt.figure(figsize=(10, 6))
sns.histplot(new_data['Present_Price'], bins=20, kde=True, color='darkviolet')
plt.xlabel('Present Price')
plt.ylabel('Frequency')
plt.title('Distribution of Present Prices')
plt.show()
```



```
In [35]: plt.figure(figsize=(10, 6))
sns.histplot(new_data['Driven_kms'], bins=20, kde=True, color='red')
plt.xlabel('Driven kilometers')
plt.ylabel('Frequency')
plt.title('Distribution of kilometers driven')
plt.show()
```



4. Preparing the data for Modelling

Since, we need to predict the price(selling price) of the car using the features of the given dataset

```
In [36]: x_data = new_data.drop(['Car_Name', 'Selling_Price'], axis=1)
y_data = new_data['Selling_Price']
```

In [37]: `x_data`

Out[37]:

	Year	Present_Price	Driven_kms	Fuel_Type_encoded	Selling_type_encoded	Transmission_e
0	2014	5.59	27000	0	0	0
1	2013	9.54	43000	1	0	0
2	2017	9.85	6900	0	0	0
3	2011	4.15	5200	0	0	0
4	2014	6.87	42450	1	0	0
...
294	2016	11.60	33988	1	0	0
295	2015	5.90	60000	0	0	0
296	2009	11.00	87934	0	0	0
297	2017	12.50	9000	1	0	0
298	2016	5.90	5464	0	0	0

299 rows × 7 columns

In [38]: `y_data`

Out[38]:

0	3.35
1	4.75
2	7.25
3	2.85
4	4.60
...	...
294	9.50
295	4.00
296	3.35
297	11.50
298	5.30

Name: Selling_Price, Length: 299, dtype: float64

In [39]: `x_train, x_test, y_train, y_test = train_test_split(x_data, y_data, test_size=`

5. Training the Models using the above data

(A) -> (i) Training using Linear Regression Model

In [40]:

```
model_1 = LinearRegression()
model_1.fit(x_train, y_train)
```

Out[40]: `LinearRegression()`

(ii) Predicting using Linear Regression Model

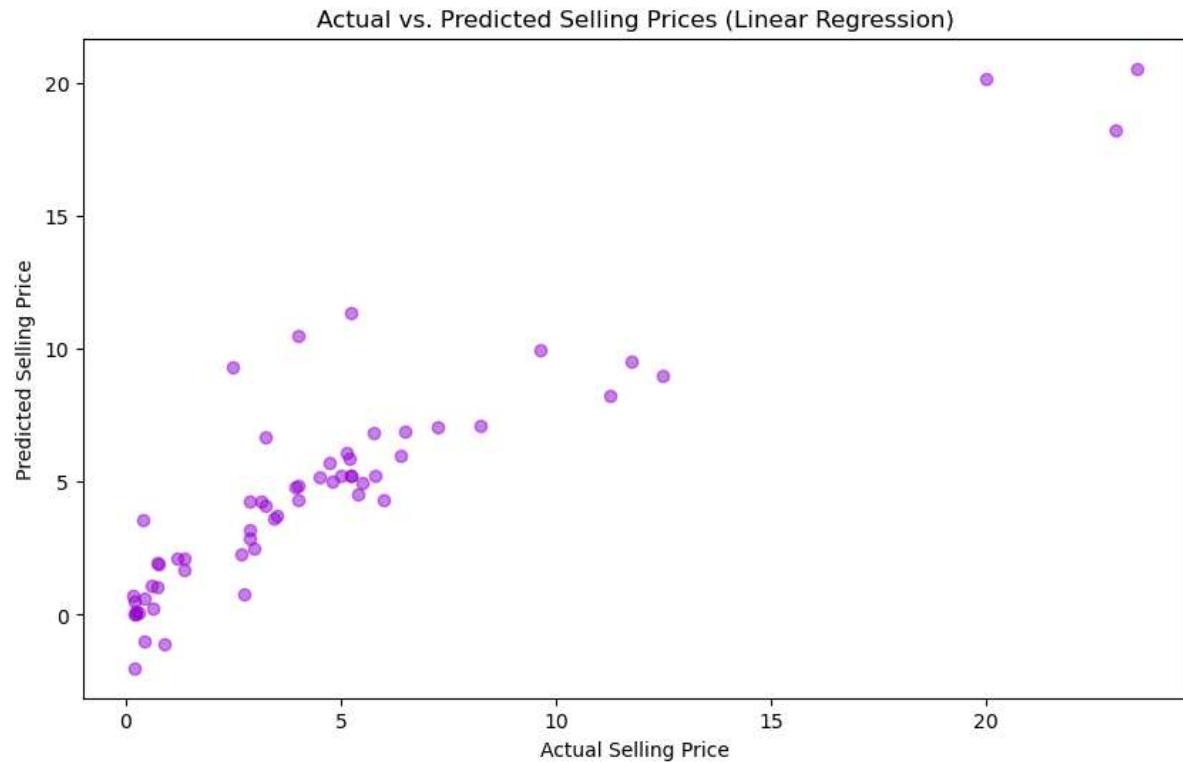
```
In [41]: y_pred_1 = model_1.predict(x_test)
```

(iii) Evaluating the performance by Linear Regression Model

```
In [42]: mse = mean_squared_error(y_test, y_pred_1)
r2 = r2_score(y_test, y_pred_1)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

Mean Square Error is : 4.028076863760065
 R-Squared score is : 0.8380235823890909

```
In [43]: plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred_1, alpha=0.5, color='darkviolet')
plt.xlabel('Actual Selling Price')
plt.ylabel('Predicted Selling Price')
plt.title('Actual vs. Predicted Selling Prices (Linear Regression)')
plt.show()
```



(B) -> (i) Training using the Random Forest Regression Model

```
In [44]: model_2 = RandomForestRegressor()
model_2.fit(x_train, y_train)
```

Out[44]: RandomForestRegressor()

(ii) Predicting using Random Forest Regression Model

```
In [45]: y_pred_2 = model_2.predict(x_test)
```

(iii) Evaluating the performance by Random Forest Regression Model

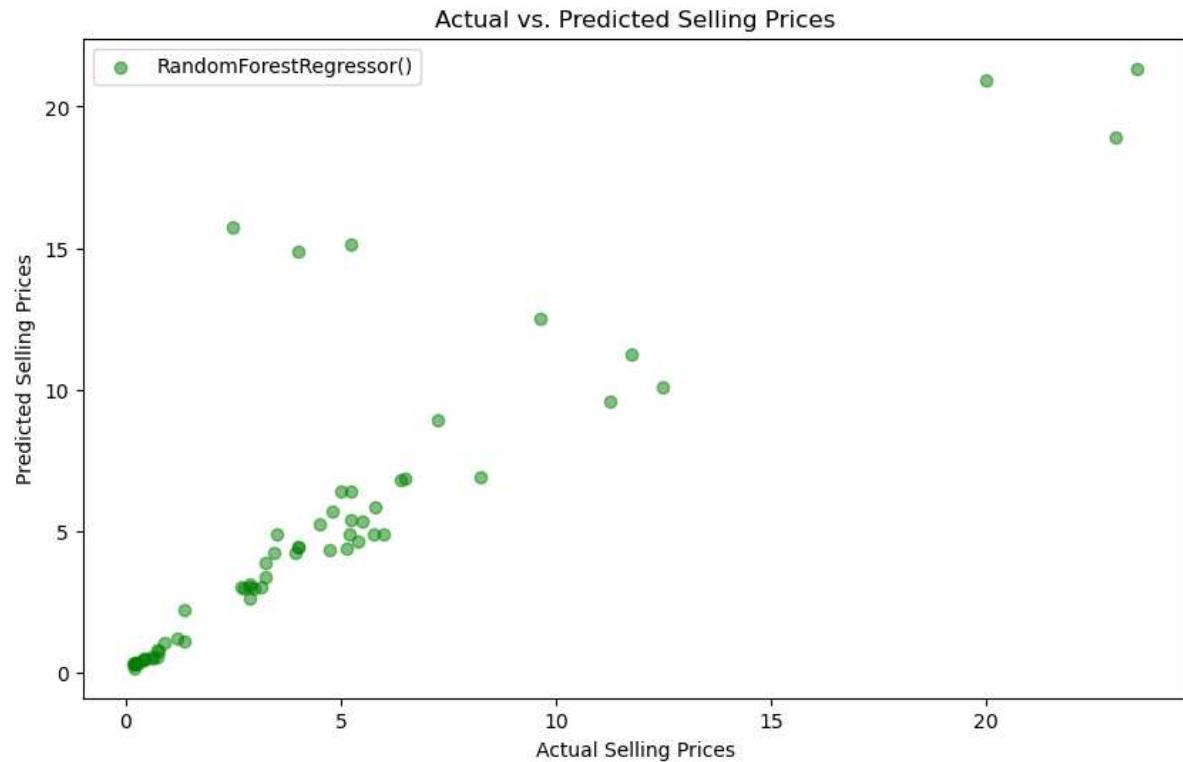
```
In [46]: mse = mean_squared_error(y_test, y_pred_2)
r2 = r2_score(y_test, y_pred_2)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

Mean Square Error is : 7.478897279166666
 R-Squared score is : 0.6992597137710552

```
In [47]: plt.figure(figsize=(10, 6))

plt.scatter(y_test, y_pred_2, label=model_2, alpha=0.5, color='green')

plt.xlabel("Actual Selling Prices")
plt.ylabel("Predicted Selling Prices")
plt.title("Actual vs. Predicted Selling Prices")
plt.legend()
plt.show()
```



(C) -> (i) Training using the Ridge Regression Model

```
In [48]: model_3 = Ridge()
model_3.fit(x_train, y_train)
```

Out[48]: Ridge()

(ii) Predicting using Ridge Regression Model

```
In [49]: y_pred_3 = model_3.predict(x_test)
```

(iii) Evaluating the performance by Ridge Regression Model

```
In [50]: mse = mean_squared_error(y_test, y_pred_3)
r2 = r2_score(y_test, y_pred_3)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

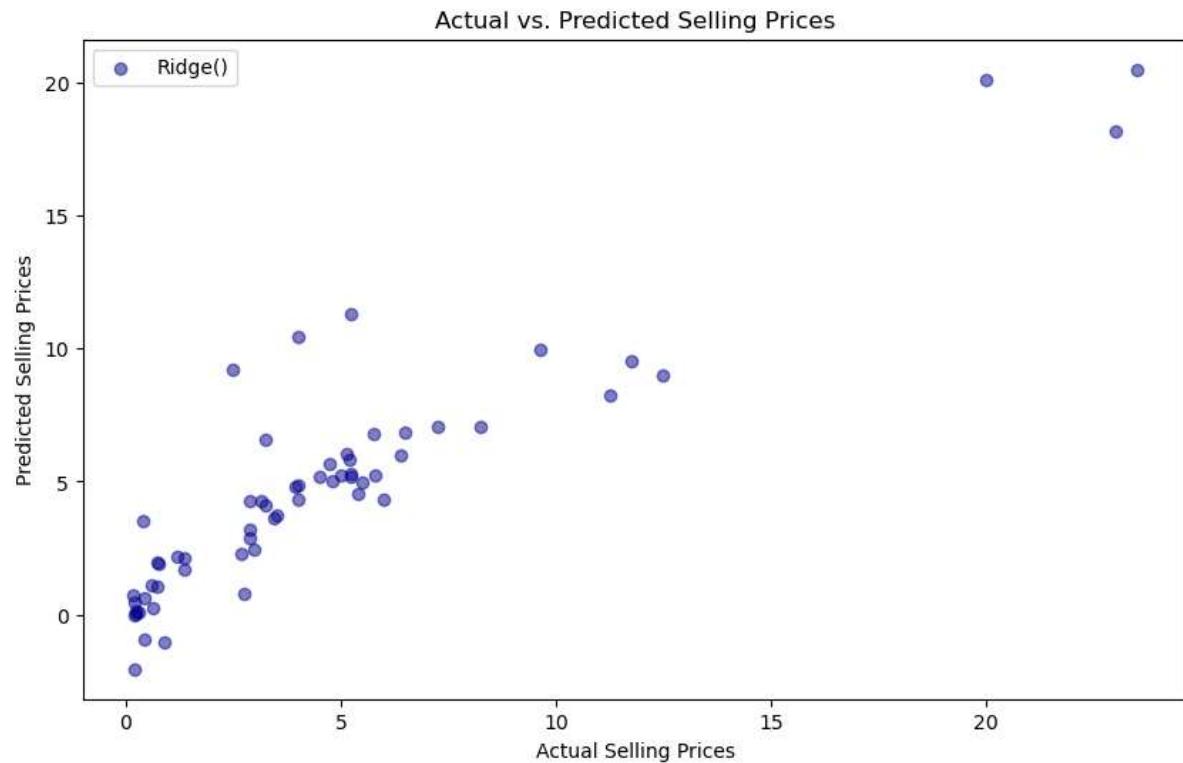
Mean Square Error is : 3.9768973298980606

R-Squared score is : 0.840081606063998

```
In [51]: plt.figure(figsize=(10, 6))

plt.scatter(y_test, y_pred_3, label=model_3, alpha=0.5, color='darkblue')

plt.xlabel("Actual Selling Prices")
plt.ylabel("Predicted Selling Prices")
plt.title("Actual vs. Predicted Selling Prices")
plt.legend()
plt.show()
```



(D) -> (i) Training using the Lasso Regression Model

```
In [52]: model_4 = Lasso()
model_4.fit(x_train, y_train)
```

```
Out[52]: Lasso()
```

(ii) Predicting using the Lasso Regression Model

```
In [53]: y_pred_4 = model_4.predict(x_test)
```

(iii) Evaluating the performance by Lasso Regression Model

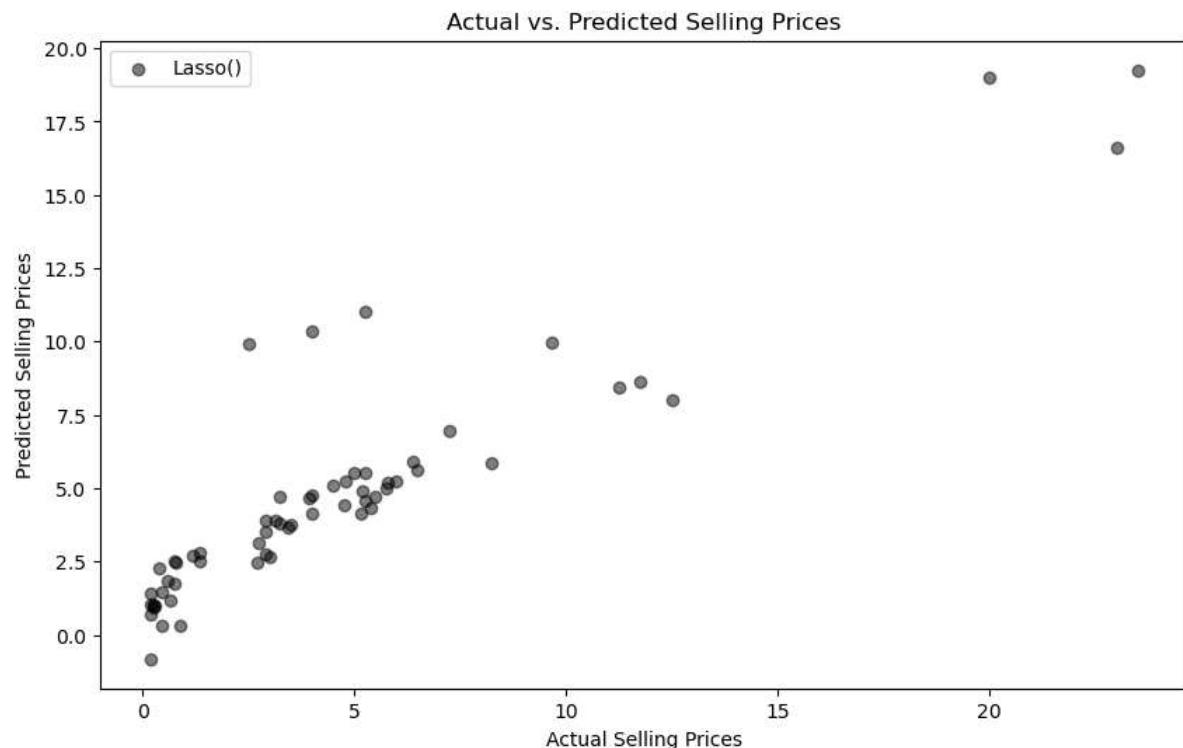
```
In [54]: mse = mean_squared_error(y_test, y_pred_4)
r2 = r2_score(y_test, y_pred_4)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

```
Mean Square Error is : 4.489972829825654
R-Squared score is : 0.8194498916620485
```

```
In [55]: plt.figure(figsize=(10, 6))

plt.scatter(y_test, y_pred_4, label=model_4, alpha=0.5, color='black')

plt.xlabel("Actual Selling Prices")
plt.ylabel("Predicted Selling Prices")
plt.title("Actual vs. Predicted Selling Prices")
plt.legend()
plt.show()
```



(E) -> (i) Training using the ElasticNet Regression Model

```
In [56]: model_5 = ElasticNet()
model_5.fit(x_train, y_train)
```

```
Out[56]: ElasticNet()
```

(ii) Predicting using the ElasticNet Regression Model

```
In [57]: y_pred_5 = model_5.predict(x_test)
```

(iii) Evaluating the performance by the ElasticNet Regression Model

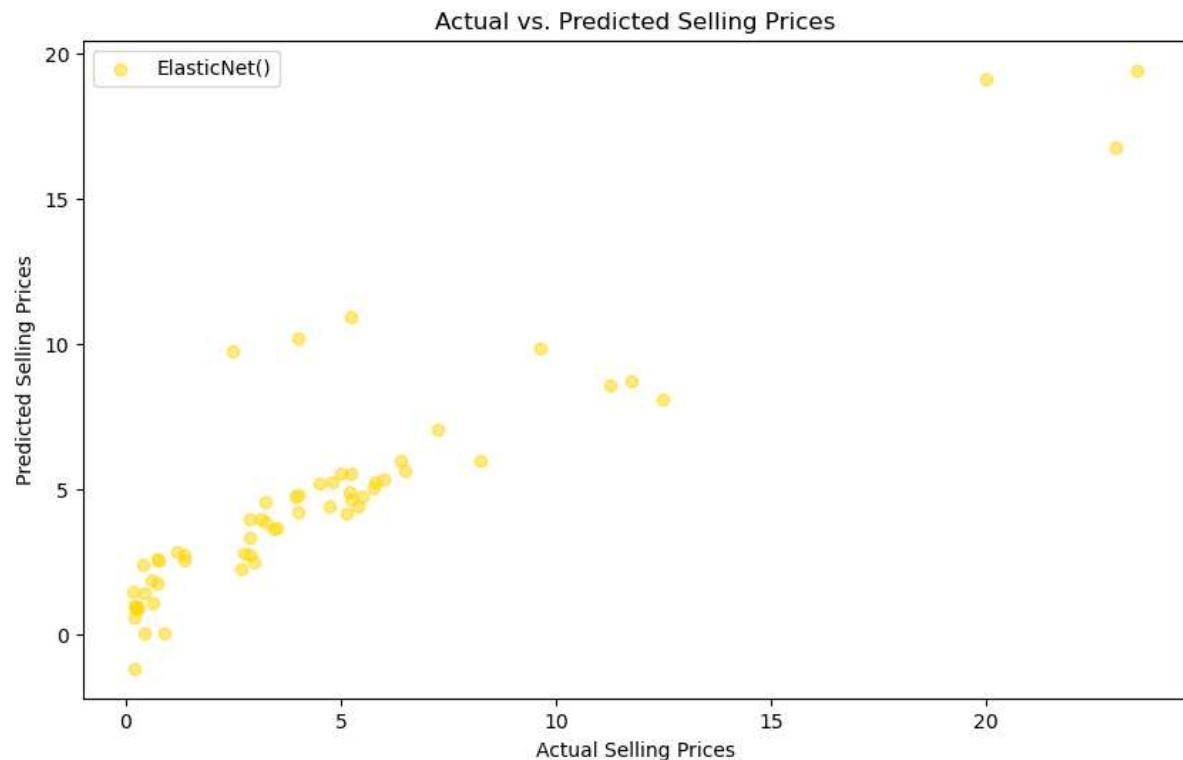
```
In [58]: mse = mean_squared_error(y_test, y_pred_5)
r2 = r2_score(y_test, y_pred_5)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

```
Mean Square Error is : 4.329718572290692
R-Squared score is : 0.825894011628066
```

```
In [59]: plt.figure(figsize=(10, 6))

plt.scatter(y_test, y_pred_5, label=model_5, alpha=0.5, color='gold')

plt.xlabel("Actual Selling Prices")
plt.ylabel("Predicted Selling Prices")
plt.title("Actual vs. Predicted Selling Prices")
plt.legend()
plt.show()
```



(F) -> (i) Training using the Gradient Boosting Regression Model

```
In [60]: model_6 = GradientBoostingRegressor()
model_6.fit(x_train, y_train)
```

```
Out[60]: GradientBoostingRegressor()
```

(ii) Predicting using the Gradient Boosting Regression Model

```
In [61]: y_pred_6 = model_6.predict(x_test)
```

(iii) Evaluating the performance of Gradient Boosting Regression Model

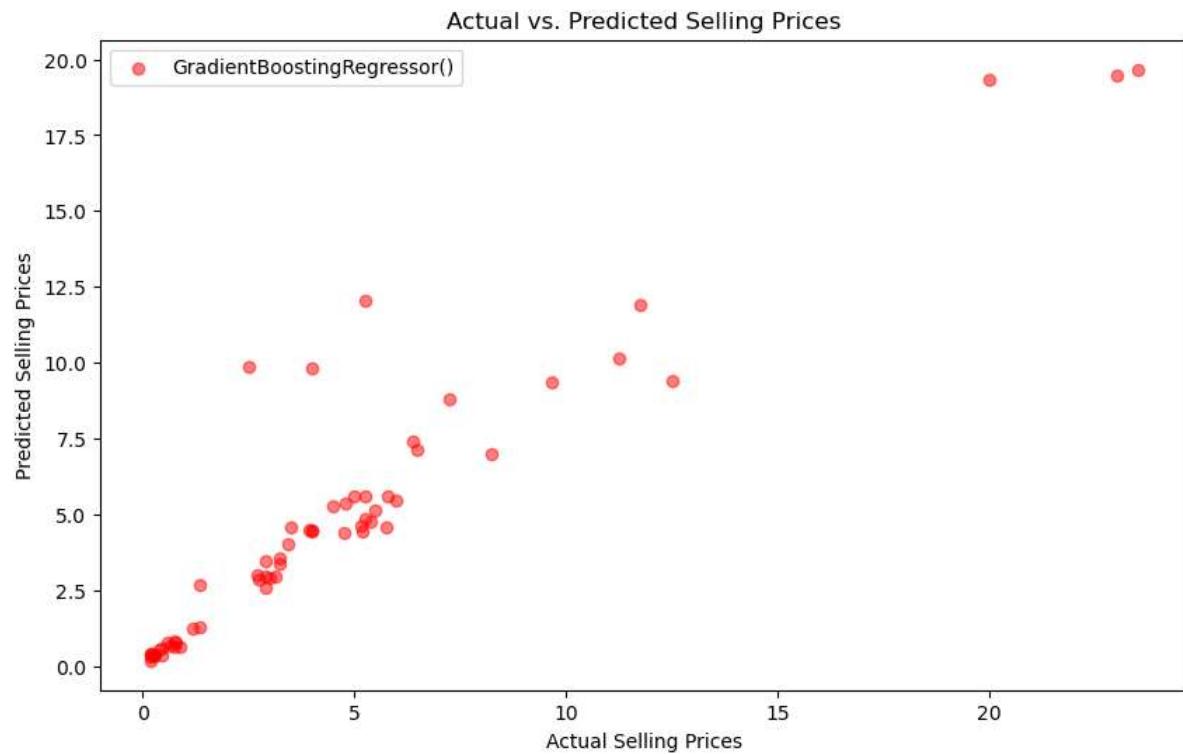
```
In [62]: mse = mean_squared_error(y_test, y_pred_6)
r2 = r2_score(y_test, y_pred_6)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

Mean Square Error is : 3.1284405404718694
R-Squared score is : 0.8741996221538454

```
In [63]: plt.figure(figsize=(10, 6))

plt.scatter(y_test, y_pred_6, label=model_6, alpha=0.5, color='red')

plt.xlabel("Actual Selling Prices")
plt.ylabel("Predicted Selling Prices")
plt.title("Actual vs. Predicted Selling Prices")
plt.legend()
plt.show()
```



(G) -> (i) Training using the XGBoost Regression Model

```
In [64]: model_7 = XGBRegressor()
model_7.fit(x_train, y_train)
```

```
Out[64]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                      colsample_bylevel=None, colsample_bynode=None,
                      colsample_bytree=None, device=None, early_stopping_rounds=None,
                      enable_categorical=False, eval_metric=None, feature_types=None,
                      gamma=None, grow_policy=None, importance_type=None,
                      interaction_constraints=None, learning_rate=None, max_bin=None,
                      max_cat_threshold=None, max_cat_to_onehot=None,
                      max_delta_step=None, max_depth=None, max_leaves=None,
                      min_child_weight=None, missing=nan, monotone_constraints=None,
                      multi_strategy=None, n_estimators=None, n_jobs=None,
                      num_parallel_tree=None, random_state=None, ...)
```

(ii) Predicting using the XGBoost Regression Model

```
In [65]: y_pred_7 = model_7.predict(x_test)
```

(iii) Evaluating the performance by XGBoost Regression Model

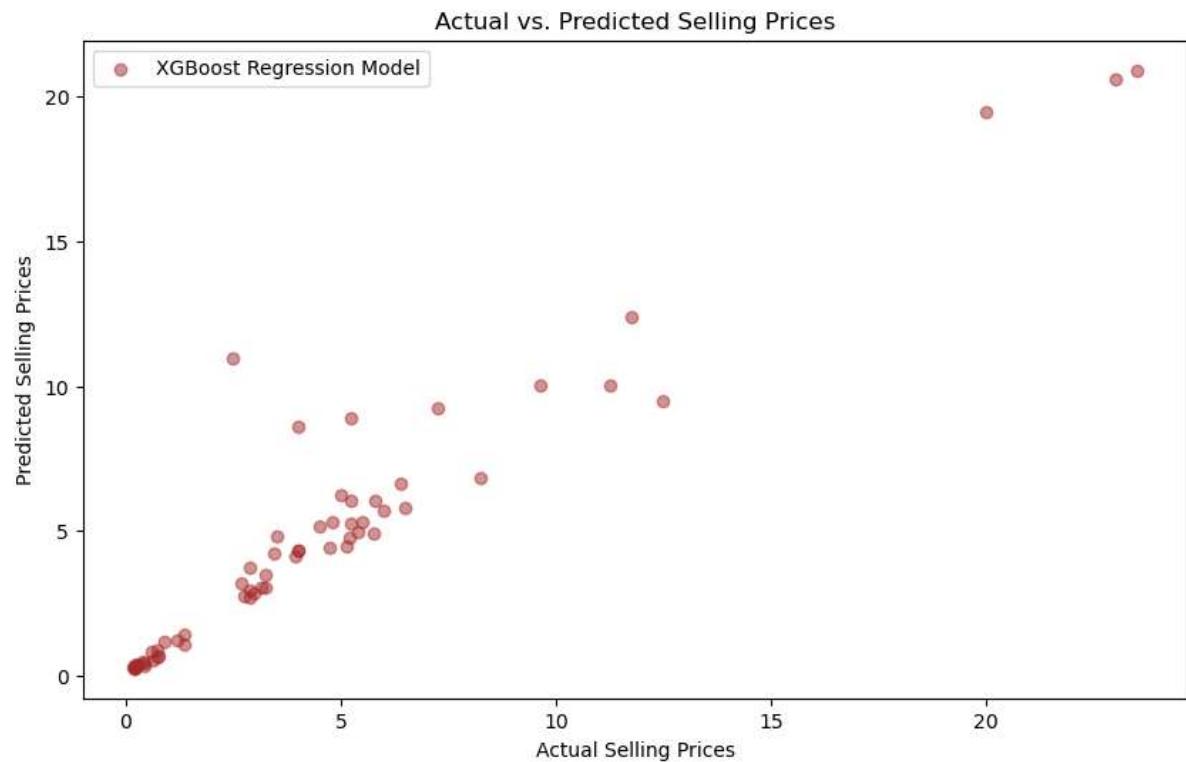
```
In [66]: mse = mean_squared_error(y_test, y_pred_7)
r2 = r2_score(y_test, y_pred_7)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
```

```
Mean Square Error is : 2.4098547850320946
R-Squared score is : 0.9030952838676375
```

```
In [67]: plt.figure(figsize=(10, 6))

plt.scatter(y_test, y_pred_7, label='XGBoost Regression Model', alpha=0.5, color='red')

plt.xlabel("Actual Selling Prices")
plt.ylabel("Predicted Selling Prices")
plt.title("Actual vs. Predicted Selling Prices")
plt.legend()
plt.show()
```



```
In [68]: print("R-Squared score of each Model:-")
print("\n Linear Regression Model: ",r2_score(y_test, y_pred_1))
print("\n Random Forest Regression Model: ",r2_score(y_test, y_pred_2))
print("\n Ridge Regression Model: ",r2_score(y_test, y_pred_3))
print("\n Lasso Regression Model: ",r2_score(y_test, y_pred_4))
print("\n ElasticNet Regression Model: ",r2_score(y_test, y_pred_5))
print("\n Gradient Boosting Regression Model: ",r2_score(y_test, y_pred_6))
print("\n XGBoost Regression Model: ",r2_score(y_test, y_pred_7))
```

R-Squared score of each Model:-

Linear Regression Model: 0.8380235823890909

Random Forest Regression Model: 0.6992597137710552

Ridge Regression Model: 0.840081606063998

Lasso Regression Model: 0.8194498916620485

ElasticNet Regression Model: 0.825894011628066

Gradient Boosting Regression Model: 0.8741996221538454

XGBoost Regression Model: 0.9030952838676375

6. Saving the Model having the best fit

In our case it's XGBoost Regression Model with R² Score of 0.9031

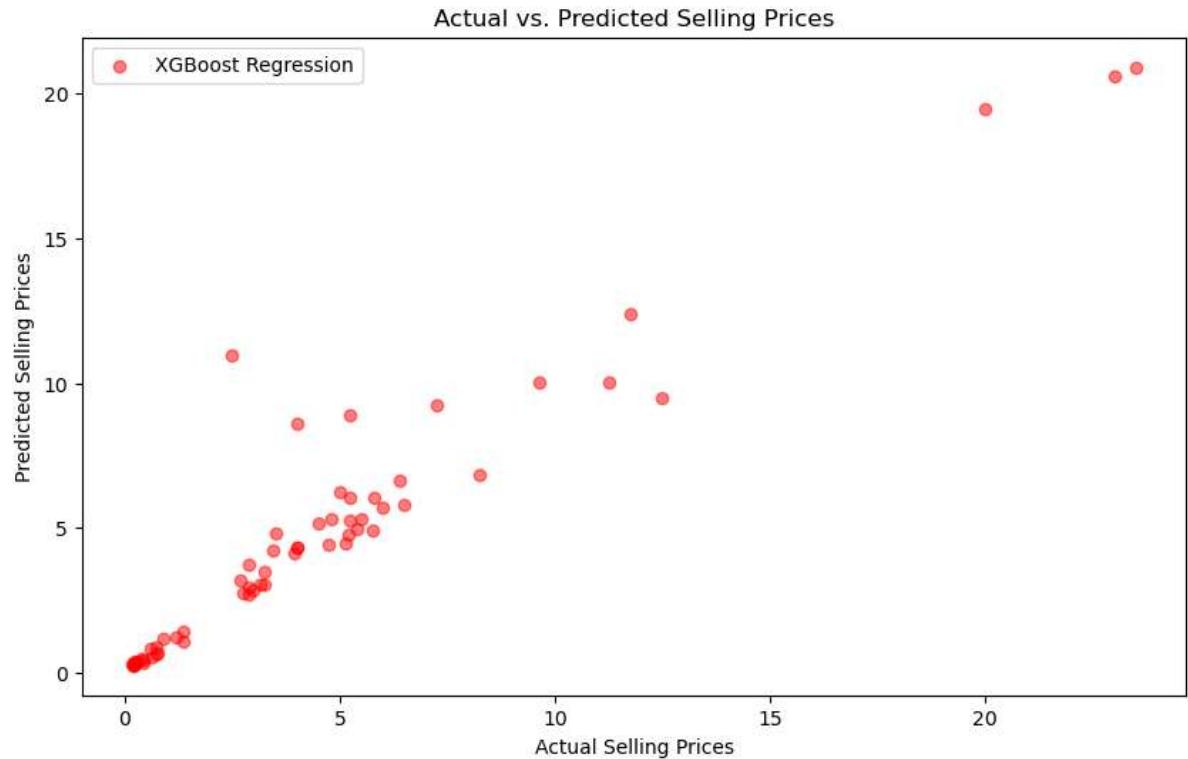
```
In [76]: final_model = XGBRegressor()
final_model.fit(x_train, y_train)
y_pred_final = final_model.predict(x_test)
mse = mean_squared_error(y_test, y_pred_final)
r2 = r2_score(y_test, y_pred_final)
print("Mean Square Error is :", mse)
print("R-Squared score is :", r2)
plt.figure(figsize=(10, 6))

plt.scatter(y_test, y_pred_final, label='XGBoost Regression', alpha=0.5, color='red')

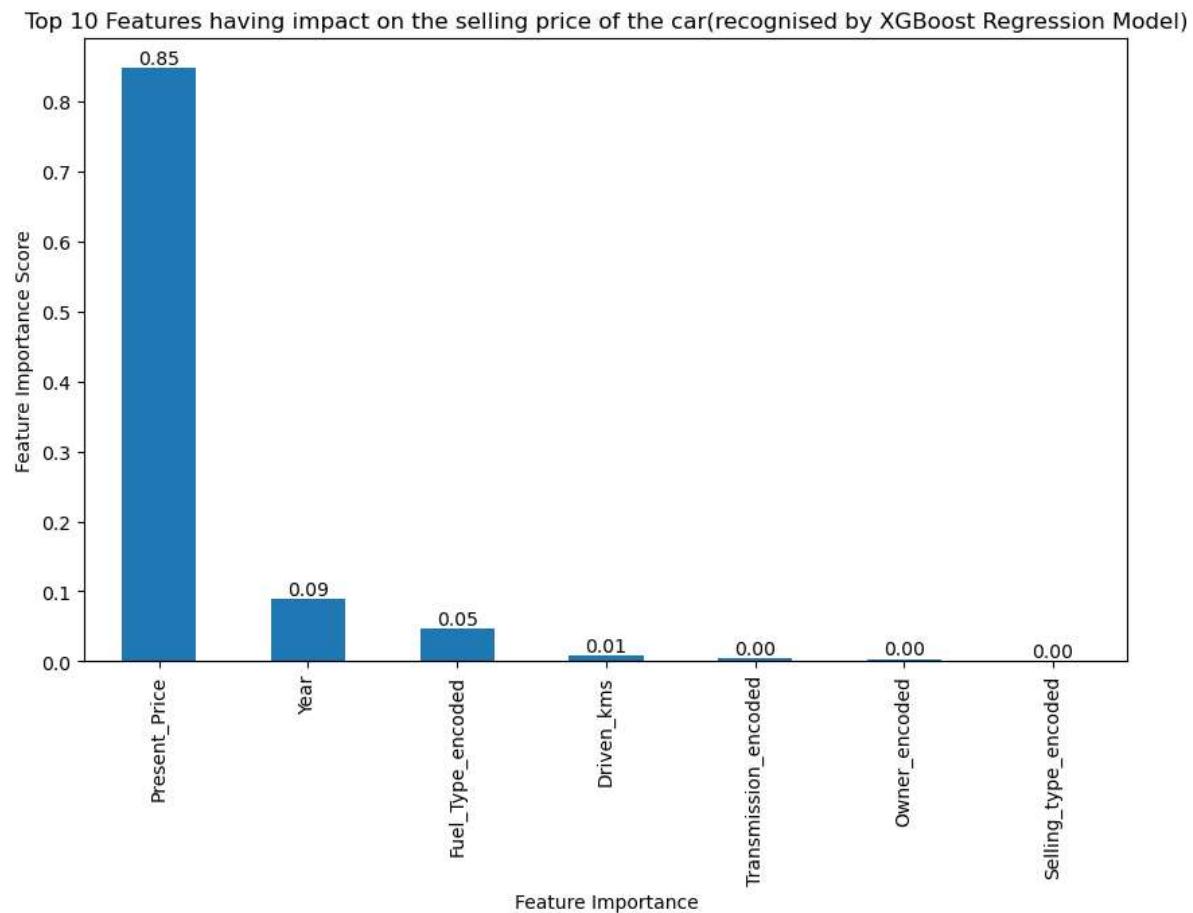
plt.xlabel("Actual Selling Prices")
plt.ylabel("Predicted Selling Prices")
plt.title("Actual vs. Predicted Selling Prices")
plt.legend()
plt.show()
```

Mean Square Error is : 2.4098547850320946

R-Squared score is : 0.9030952838676375



```
In [77]: feature_importances = pd.Series(final_model.feature_importances_, index=x_data.columns)
plt.figure(figsize=(10, 6))
top_10_features = feature_importances.nlargest(10)
feature_importances.nlargest(10).plot(kind='bar')
plt.xlabel('Feature Importance')
plt.ylabel('Feature Importance Score')
plt.title('Top 10 Features having impact on the selling price of the car(recognised by XGBoost Regression Model)')
for index, value in enumerate(top_10_features):
    plt.text(index, value, f'{value:.2f}', ha='center', va='bottom')
plt.show()
```



```
In [71]: selling_price = new_data['Selling_Price']
present_price = new_data['Present_Price']
year = new_data['Year']

scatter_selling = go.Scatter(x=year, y=selling_price, mode='markers', name='Selling Price')
scatter_present = go.Scatter(x=year, y=present_price, mode='markers', name='Present Price')

fig = go.Figure(data=[scatter_selling, scatter_present])

fig.update_layout(
    xaxis_title='Year',
    yaxis_title='Price',
    title='Trivariate Plot of Selling Price and Present Price vs. Year'
)

fig.show()
```

```
In [72]: selling_price = new_data['Selling_Price']
year = new_data['Year']

scatter_selling = go.Scatter(x=year, y=selling_price, mode='markers', name='Selling Price')

fig = go.Figure(data=[scatter_selling])

fig.update_layout(
    xaxis_title='Year',
    yaxis_title='Selling Price',
    title='Bivariate Plot of Selling Price vs. Year'
)

fig.show()
```

```
In [73]: present_price = new_data['Present_Price']
year = new_data['Year']

scatter_present = go.Scatter(x=year, y=present_price, mode='markers', name='Present Price')

fig = go.Figure(data=[scatter_present])

fig.update_layout(
    xaxis_title='Year',
    yaxis_title='Present Price',
    title='Bivariate Plot of Present Price vs. Year'
)

fig.show()
```

```
In [78]: model_filename = "final_car_prediction_model.h5"
joblib.dump(final_model, model_filename)
```

```
Out[78]: ['final_car_prediction_model.h5']
```

In []: